



Linear Regression

Mentor: Pararawendy Indarjo



Hey I'm,
Pararawendy Indarjo

I am a,

- CURRENTLY | **Senior DS at Bukalapak**
- 19 – 20 | **Data Analyst at Eureka.ai**



BSc Mathematics



Universiteit
Leiden

MSc Mathematics

Linkedin :
<https://www.linkedin.com/in/pararawendy-indarjo/>
Blog : medium.com/@pararawendy19





Outline

- Introduction to Linear Regression
- Linear Regression with Python
 - Single predictor
 - Multiple predictors
- Model diagnostic
- Model evaluation



Types of ML

RECAP

Based on target availability

1. Supervised learning

- Supervised = ground truth provided
- I.e. target is available in training data
 - (x_1, y_1) , (x_2, y_2) , (x_3, y_3) , etc..
- Objective: to correctly predict y from unseen data x
- Based on target variable type, can further be divided:
 - Regression ←
 - Classification

2. Unsupervised learning

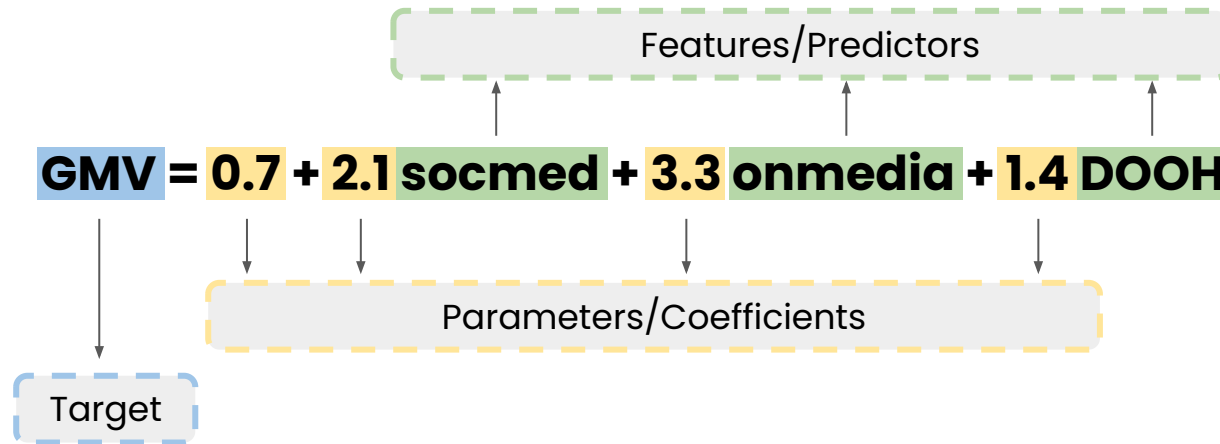
- Unsupervised = *no* ground truth provided
- I.e. target is NOT available in training data
 - (x_1) , (x_2) , (x_3) , ...
- Objective: to find hidden patterns in unlabeled data
 - Clustering
 - Dimensionality reduction



Regression Example

RECAP


Consider a model to predict sales omzet (GMV) using different advertising channels



- Regression model:
 - Predict GMV (Sales omzet)
 - Predictors are different advertising channels



First: What is a model?



RECAP



Introduction to Linear Regression

- Linear regression is the oldest statistical model, invented by Legendre in 1805
- On a high-level, linear regression is all about finding a **straight line (linear)** that fits the data most nicely
- The mathematical form


$$y = b_0 + b_1x_1 + b_2x_2 + \dots + b_nx_n$$

- y is target, x_1, x_2, \dots, x_n are **predictors/features**, $b_0, b_1, b_2, \dots, b_n$ are **parameters/coefficients** to learn
- Characteristics
 - Supervised learning (y is provided)
 - Regression model (y is quantitative/numeric)



Residual

- **Residual:** delta between the real value of target variable y and the predicted value target variable $y_{\hat{}}$
 - Predicted values = regression line



- **red points** are our data points
- **black line** is the regression line (our prediction values)


Residuals construct Loss Function

The magnitude we want to minimize when training the model

- **Residual:** delta between the real target variable y and the predicted target variable y_{hat}
 - Predicted values = regression line
- Train the model to minimize loss function: **Residual sum of squares**

$$\text{total error} = \sum_i (y_i - \hat{y}_i)^2$$

- **Note:** training == finding coefficients



- **red points** are our data points
- **black line** is the regression line (our prediction values)






Assumptions of Linear Regression

Assumption 1

- There is a straight-line relationship between the predictors and the target





- How do we check this?
 - One predictor: Scatter plot (like above)
 - Multiple predictors: Residual plot



Assumption 1 (cont'd)

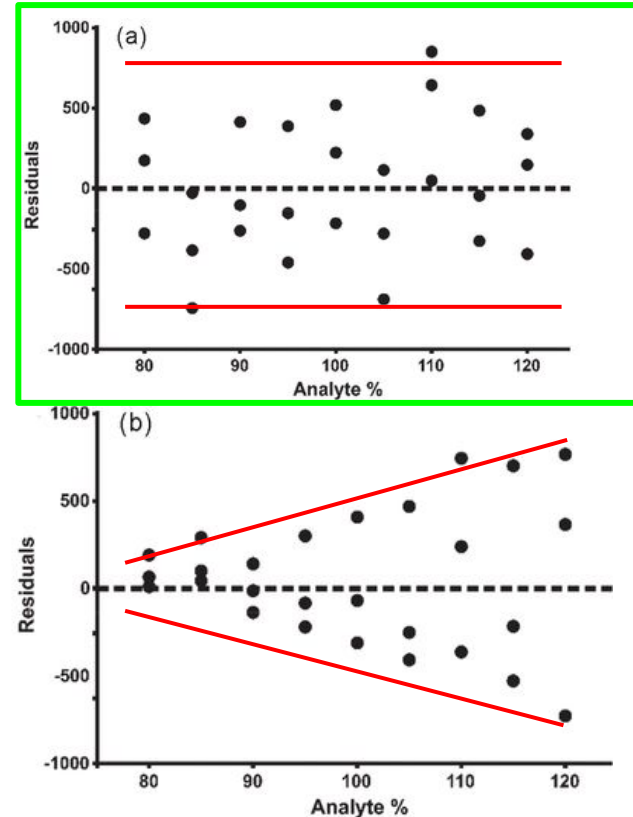
Checking linearity assumption via residual plot:

- Validated if there is NO discerning non-linear pattern in residual plot





Assumption 2

- Residuals have a constant variance




Assumption 3

- Uncorrelated residuals between different observations
- This is equivalent with independent observations




Assumption 4

- Residuals are normally distributed



- Can check using QQ-plot



Remark on Assumptions


- On real dataset, it is QUITE RARE all of those 4 assumptions are met
- It is the reason why linear regression is usually **underperformed** on real world data
- Nevertheless, **just carry on!** Real benefits of linear regression:
 - As a baseline model
 - Highly interpretable model






Essential rule in ML modelling

- We are NOT allowed to use ALL of our data to train our model
- Instead, we need to dedicate some portion of our data to become test data



Essential rule in ML modelling

- We are NOT allowed to use ALL of our data to train our model
- Instead, we need to dedicate some portion of our data to become test data



- Because we want our ML model to generalize well on new data
 - I.e. perform well on unseen data
- **Test data is meant to mimic the unseen data**



Linear Regression with Python

- **High level steps:**

- Split data: training-testing
- Prepare as numpy arrays (`X_train, y_train; X_test, y_test`)
- Define and train the model on the training data:
 - Define model: `linreg = LinearRegression()`
 - Train model: `linreg.fit(X_train, y_train)`
- Interpret and pre-evaluate the model (model diagnostic)
- Using the trained model to predict test data
 - `linreg.predict(X_test)`
- Evaluate the model
 - Using various regression metrics



Example: Simple Linear Regression

Regression with only 1 predictor

- We will use `faithful.csv` data
 - `eruptions`: Eruption time (in mins)
 - `waiting`: Waiting time to next eruption (in mins)
- We regress `eruptions` using `waiting`
- `faithful.head()`

	eruptions	waiting
0	3.600	79
1	1.800	54
2	3.333	74
3	2.283	62
4	4.533	85



Example: Simple Linear Regression

Splitting data

```
# split train test
from sklearn.model_selection import train_test_split
```

```
feature = faithful.drop(columns='eruptions')
target = faithful[['eruptions']]
```

```
ftr_train, ftr_test, tgt_train, tgt_test = train_test_split(feature,
                                                            target,
                                                            test_size=0.20,
                                                            random_state=42)
```

→ We use 80:20 split
→ For reproducibility

```
# set as numpy arrays
X_train = ftr_train.to_numpy()
y_train = tgt_train.to_numpy()
```



Example: Simple Linear Regression


Model Training and Obtained Coefficients

```
from sklearn.linear_model import LinearRegression

# define the model
simple_reg = LinearRegression()

# train the model
simple_reg.fit(X_train, y_train)
```

	feature	coefficient
0	intercept	-1.939198
1	waiting	0.076638



$$\text{eruptions} = -1.939 + 0.076 \text{ waiting}$$

- -1.939: when waiting is 0, the expected value of eruptions is -1.939
- 0.076: an increase of 1 on waiting time is associated with an increase of 0.076 on eruptions



Residual Plot

To check three assumptions of linear regression

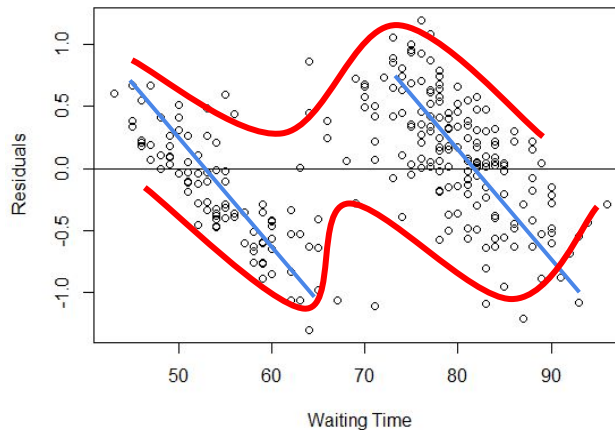
- For model with 1 predictor (feature):

- x-axis = feature value
- y-axis = residual value

```
# residual plot  
sns.scatterplot(data=df_resid, x="x_axis", y="residual")  
plt.axhline(0)  
plt.show()
```

- Assumptions to check via residual plot:

- Linear relationship ✓
- Constant variance ✗ (look at red curves)
- Independent observations ✗ (look at blue lines)




QQ-Plot

To check normality assumption of the residuals

```
# QQplot
from sklearn.preprocessing import StandardScaler

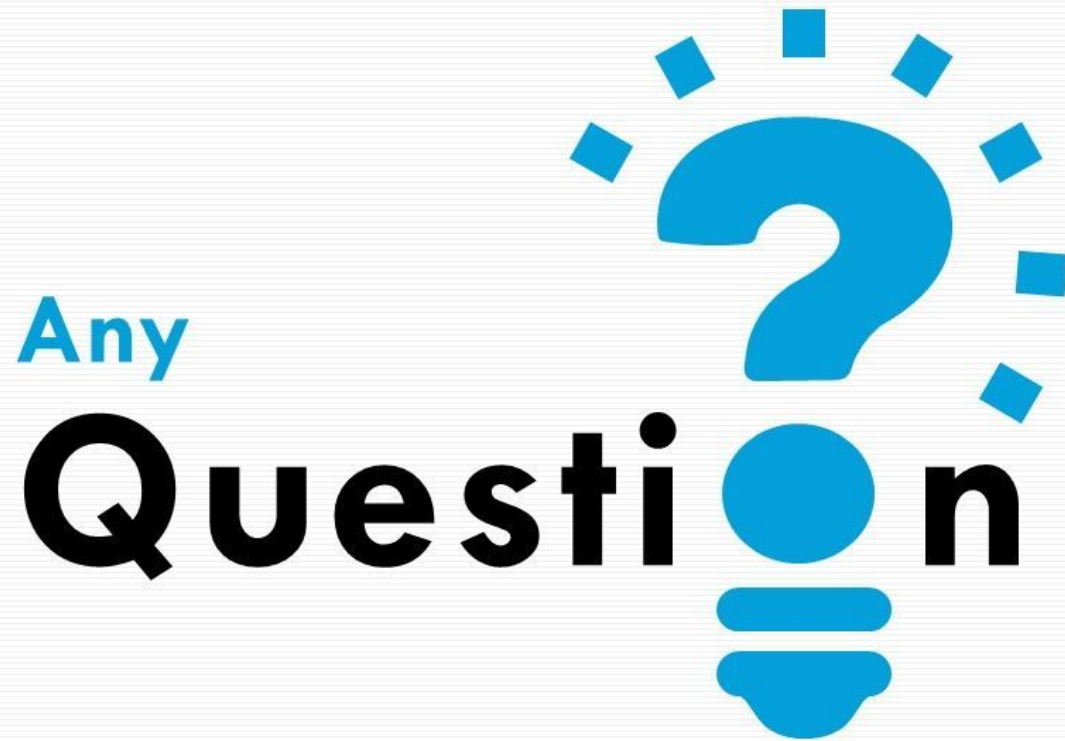
std_resid = StandardScaler().fit_transform(residual.reshape(-1,1))
std_resid = np.array([value for nested_array in std_resid for value in nested_array])

import statsmodels.api as sm
sm.qqplot(std_resid, line='45')
plt.show()
```



- QQ plot for residual
 - To check normality assumption ✓





Hands-On

- Open today's Jupyter notebook on your Google Colab!
- Make sure you have uploaded the required CSV files to your google drive
 - Remember the file path!



Multiple Linear Regression

The Data Used

- For the rest of the lecture, we will use `regression_data.csv`
- The data is about university admit probability, based on several applicant's features
 - GRE score
 - TOEFL score
 - University ranking
 - Motivation letter quality
 - Recommendation letter strength
 - GPA
 - Research experience



Modeling Flow for >1 predictor

1

Split the data

- 80 : 20 is fine

2

Multicollinearity check

- Calculate VIF score for each feature
- Correlation analysis to drop redundant feature

3

Fit the model on training data

- Define and fit `LinearRegression()`
- Only include retained features from step 2

4

Model diagnostic

- Residual plot
- QQ plot
- R2 score on training data

5

Evaluate the model on test data

- RMSE



Split the Data

Using `train_test_split()` function

- Recall, we want to predict `admit_prob`

```
# split train test
from sklearn.model_selection import train_test_split

feature = admit.drop(columns='admit_prob')
target = admit[['admit_prob']]

ftr_train, ftr_test, tgt_train, tgt_test = train_test_split(feature,
                                                             target,
                                                             test_size=0.20,
                                                             random_state=42)
```



Multicollinearity (1/2)

Variance Inflation Factor

- Multicollinearity: when two or more features/predictors are **highly correlated** with each other
- This can cause coefficients of estimates become **unreliable**
 - i.e. little change in the training data leads to completely different learned coefficients
- We can detect this by computing **Variance Inflation Factor (VIF)** for each feature

$$\text{VIF}(\hat{\beta}_j) = \frac{1}{1 - R_{X_j|X_{-j}}^2}$$

- On a high level: **VIF feature j will be high if j can be predicted using the rest of other features**. Vice versa
- $\text{VIF} == 1 \rightarrow$ No multicollinearity
- VIF between 4 and 10 \rightarrow Moderate multicollinearity
- $\text{VIF} > 10 \rightarrow$ Severe multicollinearity



Multicollinearity (2/2)

Using statsmodels library

```
from statsmodels.stats.outliers_influence import variance_inflation_factor as vif
from statsmodels.tools.tools import add_constant

X = add_constant(feature_admit_train)

vif_df = pd.DataFrame([vif(X.values, i)
                        for i in range(X.shape[1])],
                        index=X.columns).reset_index()
vif_df.columns = ['feature', 'vif_score']
vif_df = vif_df.loc[vif_df.feature!='const']
```


	feature	vif_score
1	gre_score	4.489983
2	toefl_score	3.664298
3	univ_ranking	2.572110
4	motiv_letter_strength	2.785764
5	recommendation_strength	1.977698
6	gpa	4.654540
7	research_exp	1.518065



Feature Correlation

To prevent multicollinearity


- We can draw a correlation heatmap
 - Using `sns.heatmap()`
- We found that `gre_score`, `toefl_score`, and `gpa` are highly correlated each other
 - We decide to include only `gpa` to represent these three features
 - Because it is the most correlated with the target variable
- **Note: Threshold: $abs(corr) \geq 0.8$**



Feature Correlation

To prevent multicollinearity

- We can draw a correlation heatmap
 - Using `sns.heatmap()`
- We found that `gre_score`, `toefl_score`, and `gpa` are highly correlated each other
 - We decide to include only `gpa` to represent these three features
 - Because it is the most correlated with the target variable
- **Note: Threshold: $abs(corr) \geq 0.8$**



Training the Model

Excluding gre_score and toefl_score

```
feature_admit_train = feature_admit_train.drop(columns=['gre_score', 'toefl_score'])  
feature_admit_test = feature_admit_test.drop(columns=['gre_score', 'toefl_score'])
```

```
from sklearn.linear_model import LinearRegression
```

```
# define the model
```

```
multi_reg = LinearRegression()
```

```
# train the model
```

```
X_admit_train = feature_admit_train.to_numpy()
```

```
y_admit_train = target_admit_train.to_numpy()
```

```
multi_reg.fit(X_admit_train, y_admit_train)
```



Interpreting the Obtained Model

coef_df

	feature	coefficient
0	intercept	-0.766426
1	univ_ranking	0.006984
2	motiv_letter_strength	0.004346
3	recommendation_strength	0.014776
4	gpa	0.161004
5	research_exp	0.038274

```
admit_prob = -0.763 + 0.007 univ_ranking  
             + 0.004 motiv_letter  
             + 0.015 recom + 0.161 gpa  
             + 0.038 research
```

Sample coeff interpretation:

An increase of 1 point in GPA, while the other features are kept fixed, is associated with an increase of 0.161 point in admit_prob




Residual Plot

To check three assumptions of linear regression

- For >1 predictor: residual vs predicted_target

```
# prepare dataframe
# >1 predictor --> predicted value VS residual
df_resid = pd.DataFrame({
    'x_axis': y_predict_train,
    'residual': residual
})

# residual plot
sns.scatterplot(data=df_resid, x="x_axis", y="residual")
plt.axhline(0)
plt.show()
```



- Linear relationship ✓
- Constant variance ✗
- Independent observations ✓




QQ-Plot

The same as before

```
# QQplot
from sklearn.preprocessing import StandardScaler

std_resid = StandardScaler().fit_transform(residual.reshape(-1,1))
std_resid = np.array([value for nested_array in std_resid for value in nested_array])

import statsmodels.api as sm
sm.qqplot(std_resid, line='45')
plt.show()
```



- Normality ✓
 - Slightly skewed, though



R² Score

- R² score measures portion of variability of the target variable that is successfully explained (modelled) by the features included in the model
- The higher == the better
 - Max value (not possible, though) is 100%

```
# R^2 score
from sklearn.metrics import r2_score

r2_score(y_admit_train,y_predict_train)
```

0.7986824284294713

- Interpretation: 79.86% of variability of `admit_prob` is successfully explained using all the features in the model



Hands-On

- Open today's Jupyter notebook on your Google Colab!
- Make sure you have uploaded the required CSV files to your google drive
 - Remember the file path!



Model Evaluation (1/2)

- After we are satisfied (enough) with the model, we perform model evaluation
- Essentially, we check how is the **model performance on test data**
- To do so, we will compute evaluation metric
 - Root Mean Squared Error (RMSE):
 - The standard deviation of our prediction errors with respect to the regression line
 - It is a measure of how spread out the residuals are

$$RMSE = \sqrt{\sum_{i=1}^n \frac{(\hat{y}_i - y_i)^2}{n}}$$


Model Evaluation (2/2)

```
from sklearn.metrics import mean_squared_error

print('RMSE for testing data is {}'.format(np.sqrt(mean_squared_error(y_admit_test, y_predict_test))))
```

RMSE for testing data is 0.05880540869550099

- RMSE = 0.058
 - The standard deviation of prediction errors is 0.058
 - i.e. from the regression line, the residuals mostly deviate between ± 0.058





Thank you

